Training with maximum expected value-Fixed training distribution: Deployment: Common ambiguous goal layouts Distribution shift (cheese always in top-left corner). Rare proxy-distinguishing layouts | Possible incorrect (cheese away from top-left corner). goal generalization Training with minimax expected regret-Adversarial training distribution: Deployment: Distinguishing layouts with high Distribution shift regret (amplified... ... incentivizing pursuing correct Correct goal goal in these layouts (-2). generalization